

Adaptive Fireworks Model Router

Accurate answers with fewer paid inference tokens.



Every task should use only as much model as it needs.

Most requests do not need the most expensive model. TokenPilot makes that decision locally, then escalates only when the task requires more capability.

The system balances three goals

01 Preserve answer quality

02 Minimize paid tokens

03 Stay compatible with the judging container

A local decision controls the paid path.

The router uses deterministic signals, confidence scoring, caching, and controlled escalation.

1

Classify

Recognize arithmetic, coding, proof, security, and systems signals.

2

Answer

Use a deterministic local path or the cheapest allowed Fireworks tier.

3

Judge

Escalate when confidence is below the configured threshold.

4

Cache

Reuse completed work and preserve stable prompt prefixes.

The benchmark favored the efficient tier.

On the live 83-task labeled benchmark, the cheaper model produced the strongest measured baseline.

POLICY	ACCURACY	ANSWER TOKENS
Always cheap	97.6%	22,272
Always expensive	94.0%	28,234
Current router	96.4%	22,551

Measured answer tokens; judge tokens are excluded from the submission path.

Ready for the standardized container workflow.

TokenPilot is packaged for repeatable evaluation and runtime-provided credentials.

INPUT

/input/tasks.json

ROUTER

Local decision
+ Fireworks fallback

OUTPUT

/output/results.json

Runtime configuration: FIREWORKS_API_KEY · FIREWORKS_BASE_URL · ALLOWED_MODELS

TokenPilot: Adaptive Fireworks Model Router